

Exokernel: An Operating System Architecture for Application-Level Resource Management

Paper Critique

Summary

This paper proposes another approach to operating system design that improves both efficiency of hardware resources utilization and flexibility: the exokernel architecture. An exokernel is a small kernel that provides low-level interfaces to access to hardware resources. It helps applications directly control hardware resources by providing a low-level interface to applications. To test the idea, a prototype exokernel operating system were created. Outstanding results were achieved on both efficiency and performance according to measurements.

Key ideas

The exokernel follows these design principles: *Securely expose hardware*, *Expose allocation*, *Expose names* and *Expose revocation*. It leaves policy decisions for library operating systems to decide.

The key idea of exokernel architectural design is to separate protection from management by exporting hardware resources, not emulating them like in traditional operating systems. 3 techniques were used to securely export hardware resources: *secure binding*, *visible resource revocation* and an *abort protocol*.

- *Secure binding* separate authorization from utilization of resources, improving both flexibility and performance. This technique is implemented using hardware and software caching, and downloading code.
- *Visible resource revocation* gives library operating systems and applications more control over resource allocation and deallocation and inform them about the state of hardware resources.
- *The abort protocol* deal with library operating systems that fail to satisfy the revocation protocol, in this case the kernel breaks all secure bindings with the noncomplying libraries operating system and inform them.

Background

Traditional operating systems provide access to hardware resources through a set of abstractions such as file system, processes, inter-process communication. The UNIX architecture is an example where those abstractions is exported through a system-call interface to the kernel. Those system-call are hard-coded to the kernel thus very difficult to extend or modify. The approach of forcing applications to use fixed high-level abstractions compromises both of their performance and functionality. Microkernel systems address some of those problems by minimizing the amount of kernel abstraction but still has the downsides of having fixed abstractions. Many similar researches has been done to improve extensibility of operating systems but none achieved the same degree of abstraction as the exokernel.

Contributions

This paper proposes a design for operating systems that focus on separating protection and management. This approach greatly improve flexibility in operating systems, allowing applications and libraries to implement kernel-level abstractions like inter-processes communication, file system or virtual memory. Performance is also improved as low-level hardware interactions can be implemented more efficiently.

The Aegis exokernel and the ExOS library operating is also presented in this paper. They are used to experiment with the exokernel system design.

Critique

Although theoretically, performance could be significantly increased following the exokernel approach, however, implementing functioning library operating systems and applications based on those principles is extremely difficult as they have to conform to many standards and requires a deep understanding of hardware functionalities.

Leaving policies for applications to decide could lead to less consistency, for example there is no built-in check for memory access violation or invalid address. Requiring each applications to implement reliability checking modules leads to more inconsistencies and increases size of applications.

This paper does not state clearly what type system could use the exokernel structure effectively, for example general computing systems or more specialize ones.